

Forecasting House Prices Using Smart Regression Techniques

```
In [ ]: # 1 & 2. Upload and Load the Dataset
        from google.colab import files
        import pandas as pd

        uploaded = files.upload() # Upload 'sample_house_prices.csv'
        df = pd.read_csv("sample_house_prices.csv")
        print("Data loaded successfully!")
        df.head()
```

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving sample_house_prices.csv to sample_house_prices.csv
Data loaded successfully!

```
Out[ ]:
```

	LotArea	OverallQual	YearBuilt	TotRmsAbvGrd	GarageCars	FullBath	SalePrice
0	7732	3	1953	8	2	2	267712
1	14845	2	1980	8	1	1	450090
2	8264	5	1962	6	0	3	271880
3	9859	7	1908	6	1	1	139208
4	14225	9	1961	4	1	1	111761

```
In [ ]: # 3. Data Exploration
        print(df.info())
        print(df.describe())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   LotArea          100 non-null   int64
1   OverallQual      100 non-null   int64
2   YearBuilt        100 non-null   int64
3   TotRmsAbvGrd    100 non-null   int64
4   GarageCars       100 non-null   int64
5   FullBath         100 non-null   int64
6   SalePrice        100 non-null   int64
dtypes: int64(7)
memory usage: 5.6 KB
None

```

	LotArea	OverallQual	YearBuilt	TotRmsAbvGrd	GarageCars	\
count	100.000000	100.000000	100.000000	100.000000	100.000000	
mean	10177.920000	5.390000	1961.840000	5.470000	1.410000	
std	2760.741738	2.651072	33.081387	2.293843	1.146757	
min	5537.000000	1.000000	1900.000000	2.000000	0.000000	
25%	7741.750000	3.000000	1934.500000	4.000000	0.000000	
50%	10028.000000	5.000000	1963.500000	6.000000	1.000000	
75%	12387.750000	8.000000	1991.000000	7.000000	2.250000	
max	14893.000000	9.000000	2017.000000	9.000000	3.000000	

	FullBath	SalePrice
count	100.000000	100.000000
mean	2.120000	270794.020000
std	0.80754	135362.162819
min	1.000000	51221.000000
25%	1.000000	145913.000000
50%	2.000000	269796.000000
75%	3.000000	385745.500000
max	3.000000	495678.000000

```

In [ ]: # 4. Check for Missing Values and Duplicates
print("Missing values:\n", df.isnull().sum())
print("Duplicate rows:", df.duplicated().sum())

```

```


Missing values:
LotArea      0
OverallQual  0
YearBuilt    0
TotRmsAbvGrd 0
GarageCars   0
FullBath     0
SalePrice    0
dtype: int64
Duplicate rows: 0

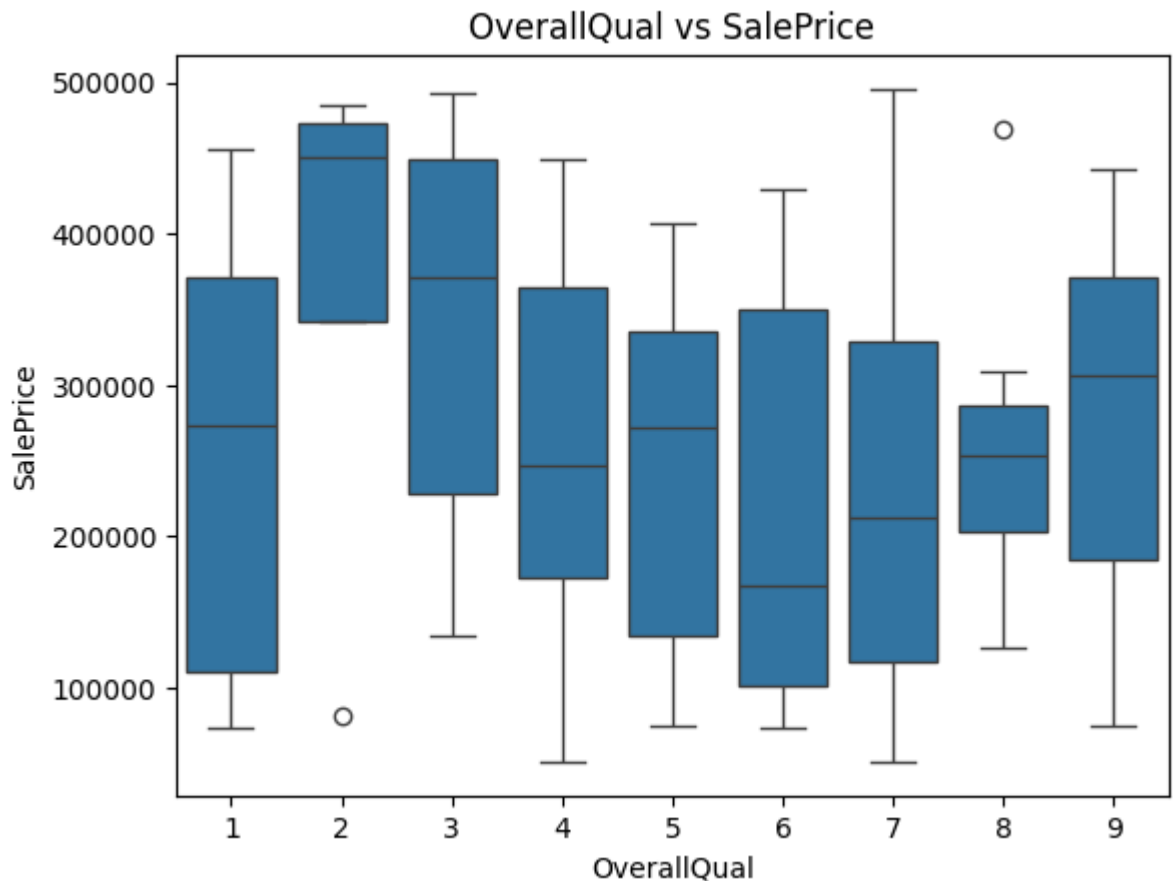
```

```
In [ ]: # 5. Visualize a Few Features
import matplotlib.pyplot as plt
import seaborn as sns

sns.histplot(df['SalePrice'], kde=True)
plt.title("SalePrice Distribution")
plt.show()

sns.boxplot(x=df['OverallQual'], y=df['SalePrice'])
plt.title("OverallQual vs SalePrice")
plt.show()
```





```
In [ ]: # 6. Identify Target and Features
```

```
X = df.drop("SalePrice", axis=1)
```

```
y = df["SalePrice"]
```

```
In [ ]: # 7 & 8. Convert Categorical Columns to Numerical and One-Hot Encode
```

```
categorical_cols = X.select_dtypes(include=['object']).columns
```

```
if not categorical_cols.empty:
```

```
    X = pd.get_dummies(X, columns=categorical_cols, drop_first=True)
```

```
In [ ]: # 9. Feature Scaling
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

```
In [ ]: # 10. Train-Test Split
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=
```

```
In [ ]: # 11. Model Building
from sklearn.linear_model import LinearRegression

model = LinearRegression()
model.fit(X_train, y_train)
```

Out []:

```
▼ LinearRegression
LinearRegression()
```

```
In [ ]: # 12. Evaluation
from sklearn.metrics import mean_squared_error, r2_score

y_pred = model.predict(X_test)
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R² Score:", r2_score(y_test, y_pred))

Mean Squared Error: 19118292320.297195
R² Score: -0.08970262913123639
```

```
In [ ]: # 13–15. Make Predictions, Convert Input to DataFrame, Predict Final Price
def predict_house_price(LotArea, OverallQual, YearBuilt, TotalBsmtSF, GrL:
                        GarageCars, GarageArea, FullBath, BedroomAbvGr, K:
    input_data = pd.DataFrame([[LotArea, OverallQual, YearBuilt, TotalBsm
                                GarageCars, GarageArea, FullBath, Bedroom
                                columns=X.columns)
    input_scaled = scaler.transform(input_data)
    prediction = model.predict(input_scaled)[0]
    return round(prediction, 2)
```

```
In [ ]: # Prediction Function (for sample_house_prices.csv)
def predict_house_price(LotArea, OverallQual, YearBuilt, TotRmsAbvGrd, Ga
    input_data = pd.DataFrame([[LotArea, OverallQual, YearBuilt, TotRmsAb
                                columns=['LotArea', 'OverallQual', 'YearBui

    input_scaled = scaler.transform(input_data)
    prediction = model.predict(input_scaled)
    return prediction[0]

# Example Prediction
predicted_price = predict_house_price(8000, 7, 2005, 6, 2, 2)
print("Predicted House Price: ${:,.2f}".format(predicted_price))

Predicted House Price: $307,317.35
```